

SOC-project

Realization Report

Baisangur Dudayev
Bachelor Electronics-ICT - Cloud & Cybersecurity

Table of Contents

TABLES AND FIGURES	3
1. INTRODUCTION	4
2. ANALYSIS	5
2.1. Tools and Technologies Used	5
3. REALISATIONS	7
3.1. Basic SIEM configuration	7
3.2. Visualization & Reporting	12
3.3. Integrating SIEM automation	17
3.4. Integrating TheHive	21
4. CONCLUSION	21
5. REFERENCES	23

TABLES AND FIGURES

Figure 1. Indices of stored logs	8
Figure 2. Extensive list of python scripts	9
Figure 3. documents count of indices before script deletes documents	10
Figure 4. documents count of indices after script deletes documents	11
Figure 5. Guardian 360 Report cover page	14
Figure 6. Guardian 360 Report table of contents	15
Figure 7. Guardian 360 Report Intro and risk table	16
Figure 8. Guardian 360 Report Donut chart	16
Figure 9. Guardian 360 Report stacked barchart	16
Figure 10. Issues Guardian 360 Issues pages	16
Figure 11. Script used to generate report, 700 lines of code	17
Figure 12. Script template used to generate report, 1100 lines of code	17
Figure 13. Running n8n container with nodes to and act on retrieve mails	18
Figure 14. Nn8n node code to take action on email's contents	19
Figure 15. Google app script to delete read mails	19
Figure 16. Playbook to run n8n and NGINX container	20
Figure 17. Updated Vocal Notify workflow to make calls based on logs	20
Figure 18. TheHive running	21

1. Introduction

Netcure provides security services to its clients through continuous monitoring and alerting. To support this, it required a Security Operations Center (SOC), which plays a key role in collecting, analysing, and responding to security events from various platforms.

However, alerts and notifications were previously spread across different tools and systems, making it difficult to get a clear overview or generate consistent reports. This created a need for a more centralized and automated approach to improve visibility and efficiency.

The goal of the internship was to help solve this by developing a system that brings together alerts and logs from different sources into one environment. This includes integrating tools such as email, APIs, ELK stack, and visual dashboards like Grafana and Kibana. The SOC not only supports client environments but also helps protect the internal infrastructure of Netcure itself.

By improving how alerts are collected, visualized, and reported, the project directly supports Netcure's core business and strengthens its security operations.



The timeline below reflects the total time I personally spent on each major topic during the internship:

- **Basic SIEM configuration 6 WEEKS**
 - Basic configuration. Collecting logs, filtering and aggregating them.
 - Setting up testing environments for my team using WSL to run Ansible commands needed for configuration and deployment tasks
- **Reporting/Visualization 4 WEEKS**
 - Integrating an automatic Guardian360 report
- **Integrating SOC automations 2,5 WEEKS**
 - Integrating n8n workflows into SIEM for alerting
 - Engineering a time based script to trim excessive documents for indices.
- **TheHive & MISP 0.5 WEEK**
 - Setting up TheHive and Cassandra, and ensuring they work smoothly with the existing containers

This document covers the following parts:

- An analysis of tools and decisions made
- A description of the realisations, covering each implemented stage
- A conclusion reflecting on the progress and what lies ahead

2. Analysis

At the beginning of the project, the employee from the internship company who was involved in the SOC project mentioned that he was very open to suggestions regarding the technologies to be used. The main requirements were that the tools had to be free and preferably popular, as widely used tools usually have better documentation and are generally considered more secure.

The internship team consisted of four interns and the employee. So even though there was a lot of freedom in choosing the tools, the final decisions were not always made by a single intern. Additionally, two Thomas More Mechelen interns started their internship one week earlier, which meant that some of the technologies had already been chosen by the time the rest of the team joined.

2.1. Tools and Technologies Used

WSL (Windows Subsystem for Linux) + Visual Studio Code

- used to create Linux environments directly on Windows machines.
- an easy and fast way approach was needed to run Ansible playbook commands on Windows, which is not possible under normal circumstances, so I chose to learn and install WSL.

Ansible

- A configuration management tool that was used to automate the provisioning and setup of infrastructure components.
- There was a need to automate the infrastructure setup using Infrastructure as Code (IaC) instead of configuring everything manually. This approach allowed me to first test the infrastructure locally and then deploy it to the production server simply by running the Ansible playbook.

Docker

- A tool that that was used to containerize and orchestrate services such as n8n, nginx, etc. ...
- The n8n installation documentation used docker to set up n8n and nginx. I followed this flow and did the same instead of running the services directly on the virtual machine of the live server.

Elasticsearch, Logstash, Kibana & (ELK Stack) Grafana

- Elasticsearch stores log data, logstash filters logs before sending them to Elasticsearch, Kibana and Grafana visualize data via dashboards and graphs.
- The ELK stack is a free, open-source, and time-tested solution. It offers extensive technical documentation, community guides, and tutorials, making it an obvious choice.

n8n and AI agents (Hugging Face):

- n8n is used to automate workflows through a visual interface, allowing users to connect different services and define logic without having to write traditional code. It supports custom logic through JavaScript when needed. AI agents can be combined with Large Language Models to perform specific tasks based on the tools and data made available to them.
- An AI specialist introduced various tools relevant to the project. To explore their potential use, a short course on AI agents (Hugging Face) was followed as part of the research phase.

Nginx

- Nginx acts as an intermediate between services like e.g. n8n. This makes certain communications possible and more secure.
- This was needed this to make communication with n8n possible and secure. Before implementing this, I had to read up on some documentation regarding nginx.

OpenSSL

- A command-line tool that was needed to generate self-signed certificates for internal HTTPS usage.
- Self-signed certificates are needed to make HTTPS (secure) connections between certain services possible.

Jinja

- A templating engine that I used to fill in dynamic fields using a script.
- This tool was also used by another script that made the Cyberalarm report. I tried it out and it worked as expected.

Python

- Python is a widely used programming language known for its readability, versatility, and strong support for scripting and automation tasks.
- It was used during the internship for tasks such as automating infrastructure setup, parsing data, and interacting with APIs,

SCP (Secure Copy Protocol)

- A protocol used for transferring files securely between machines.
- SCP provided me with a quick way to copy from the live server to the desktop computer I was working on.

VPN

- A Virtual Private Network is used to securely connect remote systems to the project network.
- The desktop computer I was working on needed a VPN to connect to the network where the live server hosting the SIEM was running.

3. Realisations

3.1. Basic SIEM configuration

During the initial phase of the internship, the SIEM environment was successfully deployed on the production server. Real-time data was collected at five-minute intervals. The other interns successfully visualized the logs using Grafana. The setup process was automated using Infrastructure as Code (IaC) principles through Ansible playbooks, which were developed and configured for services such as Elasticsearch and Kibana. (Elastic, 2025) (elasticsearch, 2025) (Stringer, 2025)

To facilitate development and testing, Windows Subsystem for Linux (WSL) was implemented on all project laptops. This enabled the team to run and test Ansible playbooks directly on Windows systems, allowing for local validation before deployment to the live environment.

Log collection and preprocessing workflows were configured in collaboration with another intern. These workflows ensured that log data was correctly parsed, cleaned, and forwarded to the appropriate services. To support this, approximately thirty scripts were developed by me to retrieve log data from various cybersecurity platforms used by client systems, including Guardian360, Cyberalarm, and WatchGuard. These scripts automated the collection of data through scheduled API calls and were tailored to the specific structure and requirements of each service.

Although some scripts required more complex logic, most shared a reusable structure. This modular approach improved development speed and allowed for easier maintenance and adaptation to changing API endpoints or data formats.

Support was also provided by me to team members from the Thomas More Mechelen campus by assisting with troubleshooting tasks and offering guidance on Linux fundamentals and command-line usage.

Additionally, a solution was developed to manage storage within Elasticsearch by automatically deleting outdated documents based on age and index size. This was achieved through a custom Python script that interfaced with Elasticsearch's REST API. The script was tested using sample indices and dry-run functionality to ensure correct behaviour before production deployment.

Available indices

☒ Show hidden indices ☒ Only show crawler indices

Index name	Index health	Docs count	Ingestion name	Ingestion method	Ingestion status	Actions
[REDACTED]	● yellow	77		API	Connected	
[REDACTED]	● yellow	107		API	Connected	
[REDACTED]	● yellow	5		API	Connected	
[REDACTED]	● yellow	1		API	Connected	
[REDACTED]	● yellow	1		API	Connected	
[REDACTED]	● yellow	24281		API	Connected	
[REDACTED]	● yellow	216		API	Connected	
[REDACTED]	● yellow	308		API	Connected	
[REDACTED]	● yellow	3323		API	Connected	
[REDACTED]	● yellow	20		API	Connected	
[REDACTED]	● yellow	1		API	Connected	
[REDACTED]	● yellow	31		API	Connected	
[REDACTED]	● yellow	18		API	Connected	
[REDACTED]	● yellow	60		API	Connected	
[REDACTED]	● yellow	66		API	Connected	
[REDACTED]	● yellow	4		API	Connected	
[REDACTED]	● yellow	11		API	Connected	
[REDACTED]	● yellow	2		API	Connected	
[REDACTED]	● yellow	3654		API	Connected	
[REDACTED]	● yellow	159		API	Connected	

< 1 2 3 4 5 ... 17 >

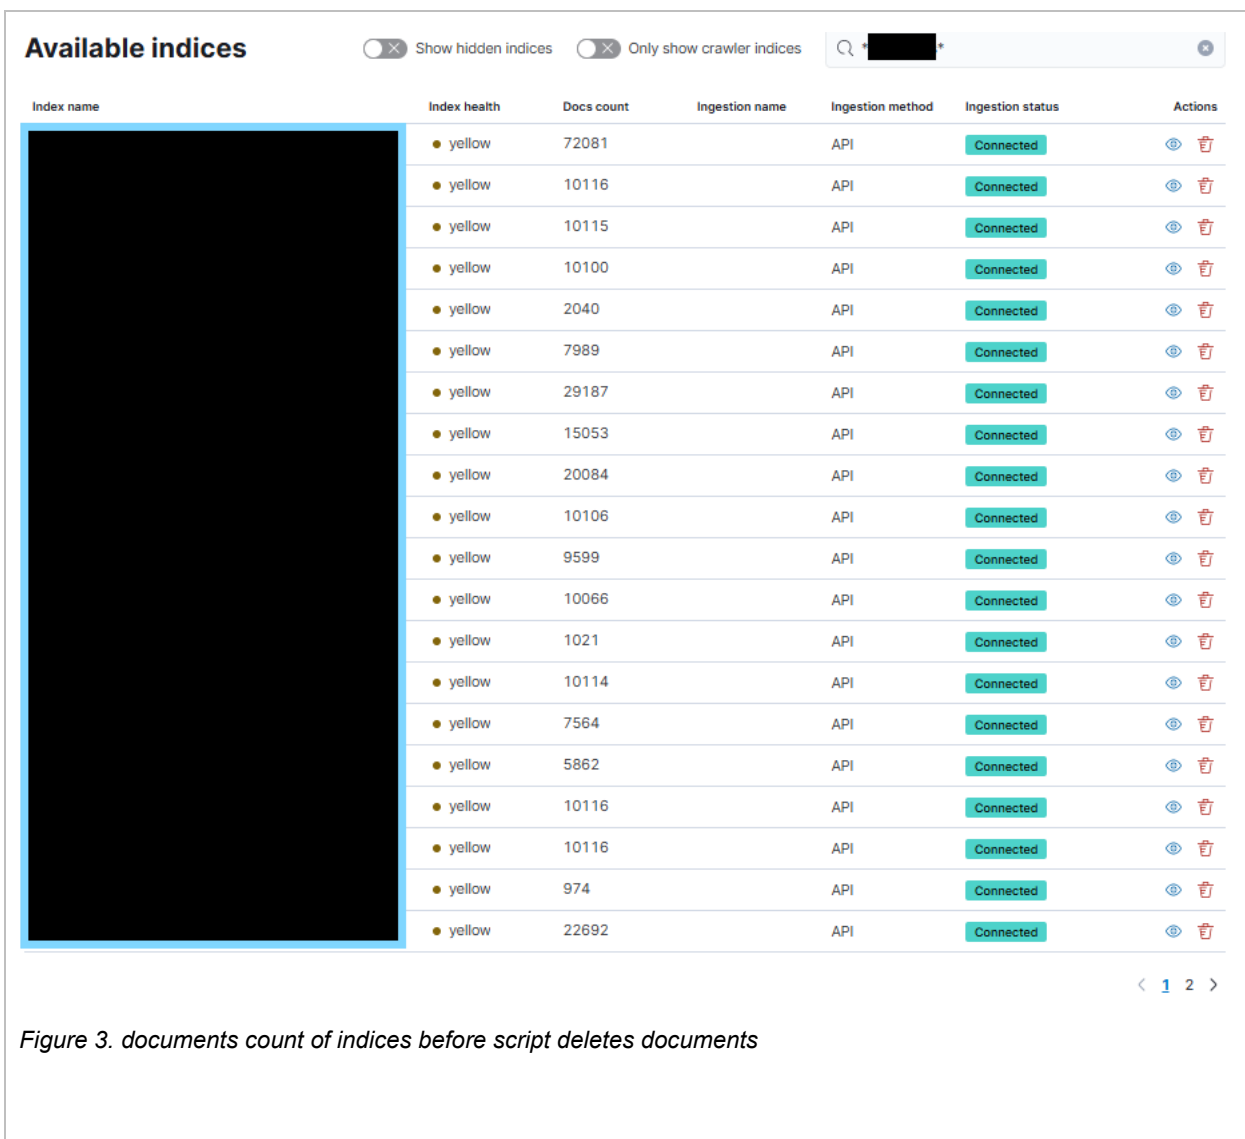
Figure 1. Indices of stored logs

The screenshot displays a code editor interface. On the left, the 'EXPLORER' sidebar shows a project structure with a folder named 'SCRIPTS SIEM'. Inside this folder, there are numerous Python files, including 'get_a_list_of_hacker_alerts.py', 'get_a_list_of_assets.py', 'get_a_list_of_companies.py', 'get_a_list_of_hacker_alert_appliances.py', 'get_a_list_of_hacker_alert_exclusions.py', 'get_a_list_of_hacker_alerts.py', 'operations_on_issues.py', 'get_a_list_of_probes.py', 'get_a_list_of_schedules.py', 'get_a_list_of_scanner_exclusions.py', 'get_a_list_of_scanner_platforms.py', 'get_live_asset_counter NEW (company fil...', 'NOT WORKING generate_usage_report.py', 'NOT WORKING retrieve_historical_trends...', 'retrieve_all_unique_issues NEW.py', 'retrieve_historical_trends_by_risk CANT F...', 'retrieve_historical_trends_by_risk.py', 'retrieve_metrics (company filter).py', 'retrieve_risks_per_detection_date (compa...', 'retrieve_risks_per_scanner (company filte...', 'retrieve_risks_per_scanobject (company f...', 'AI prompt.txt', 'api_token.py', 'compare.py', 'fieldnames to be split.txt', 'GD_API_token.txt', 'list of scripts 01-04.txt', 'retrieve_risks_per_scanobject copy.py', 'guardian360_copy_scripts_to_production', 'logstash', 'backups', 'bais_log_home_dir.conf', 'logstashB_bryan.conf', 'watchguard', 'backups', 'retrieve_a_list_of_incidents.py', and 'retrieve_filtered_ca_alarms_test.py'. The file 'get_a_list_of_hacker_alerts.py' is selected and highlighted.

The main code editor on the right shows the content of 'get_a_list_of_hacker_alerts.py'. The code is written in Python and includes the following elements:

- Imports: `import requests`, `import logging`, and `import json`.
- Comments: `# Define the API key and token file path (close to each other for easy copy-pasting)`, `# Base URL and endpoint`, `# File Paths for Logging and Output`, `# Setup logging configuration to log errors to the specified file`, `# Function to read the API token from the file`, `# Define headers, including the API key and token for authorization`, `# Function to retrieve the list of hacker alerts`, and `# Check if the response is successful (status code 200)`.
- Variables: `api_key = "REDACTED"`, `token_file_path = "REDACTED"`, `base_url = "REDACTED"`, `endpoint = "REDACTED"`, `error_log_file = "REDACTED"`, `output_file = "REDACTED"`, `api_token = read_api_token(token_file_path)`, and `headers = { "Authorization": f"Bearer {api_token}", "Api-Key": api_key, "Content-Type": "application/json" }`.
- Functions: `def read_api_token(file_path):` and `def retrieve_hacker_alerts():`.
- Code Logic: The `read_api_token` function attempts to read the API token from a file. The `retrieve_hacker_alerts` function defines parameters for pagination and filter by acknowledged status, sends a GET request to retrieve the list of hacker alerts, and checks if the response is successful.

Figure 2. Extensive list of python scripts



Elasticsearch indices

Default settings

Create a new index

Enterprise Search has not been configured

The Elastic web crawler is not available without Enterprise Search.

Review setup guide

Available indices

Show hidden indices

Only show crawler indices

Search

Index name	Index health	Docs count	Ingestion name	Ingestion method	Ingestion status	Actions
	yellow	2001		API	Connected	
	yellow	2000		API	Connected	
	yellow	2000		API	Connected	
	yellow	2000		API	Connected	
	yellow	2000		API	Connected	
	yellow	2000		API	Connected	
	yellow	2000		API	Connected	
	yellow	2000		API	Connected	
	yellow	2000		API	Connected	
	yellow	2000		API	Connected	
	yellow	2000		API	Connected	
	yellow	2000		API	Connected	
	yellow	2000		API	Connected	
	yellow	2000		API	Connected	
	yellow	22		API	Connected	
	yellow	2000		API	Connected	
	yellow	2000		API	Connected	
	yellow	2000		API	Connected	
	yellow	2000		API	Connected	
	yellow	0		API	Connected	
	yellow	2001		API	Connected	

< 1 2 >

Figure 4. documents count of indices after script deletes documents

Elasticsearch indices

[Default settings](#)[Create a new index](#)

⚠ Enterprise Search has not been configured

The Elastic web crawler is not available without Enterprise Search.

[Review setup guide](#)

Available indices



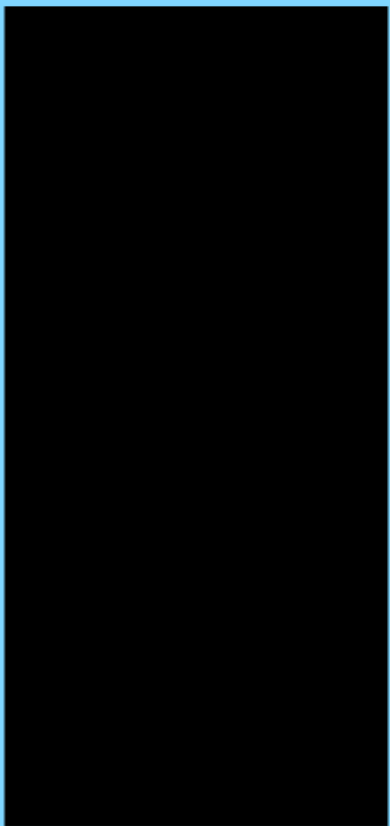










































Show hidden indices



Only show crawler indices



Search

Index name	Index health	Docs count	Ingestion name	Ingestion method	Ingestion status	Actions
	● yellow	2001		API	Connected	 
	● yellow	2000		API	Connected	 
	● yellow	2000		API	Connected	 
	● yellow	2000		API	Connected	 
	● yellow	2000		API	Connected	 
	● yellow	2000		API	Connected	 
	● yellow	2000		API	Connected	 
	● yellow	2000		API	Connected	 
	● yellow	2000		API	Connected	 
	● yellow	2000		API	Connected	 
	● yellow	2000		API	Connected	 
	● yellow	2000		API	Connected	 
	● yellow	2000		API	Connected	 
	● yellow	2000		API	Connected	 
	● yellow	22		API	Connected	 
	● yellow	2000		API	Connected	 
	● yellow	2000		API	Connected	 
	● yellow	2000		API	Connected	 
	● yellow	2000		API	Connected	 
	● yellow	0		API	Connected	 
	● yellow	2001		API	Connected	 

< 1 2 >

3.2. Visualization & Reporting

During this phase, I focused on automating the generation of a security report based on logs retrieved from Guardian360's vulnerability scans. These logs came in CSV format, consisting of 2 SCV files. I processed them into both interactive and printable reports. I developed an HTML template for the report layout and wrote a Python script that takes data from the log files and fills in the template automatically. In addition to HTML, I incorporated JavaScript for interactivity and CSS for styling. I added charts to show key metrics clearly, making the report easier to read and understand.

in the end, I had a dynamic, interactive HTML report with:

- A table of contents
- Page numbering
- A styled cover page
- Automatically populated issue descriptions, solutions, technical data, and severity levels

A PDF version of the report is also generated besides the HTML version; it kept the look of the HTML version and made it more archive friendly. I tested the reports regularly to make sure all the data was accurate and displayed correctly. Additionally, I implemented multilingual support by adding translation functionality for both static and dynamic content. While predefined translations were used for static sections such as the introduction, the dynamic fields including issue descriptions, solutions, and technical details were translated using the DeepL API. This allows the report to be generated in Dutch, French, or English, depending on the target audience. (Deepl, 2025). The DeepL functionality was later scrapped due to concerns about inaccurate translations of certain technical terms, although the code itself was fully functional.

Two CSV files are used to generate the reports. The first is the main CSV file containing the core scan data, while the second serves as a mapping file to link each company's scanner platform to the correct log entries. Although the main CSV does not directly specify the scanner platform used, it does include the scan object field. This field is also present in the second CSV, which makes it possible to establish the correct connections between platforms and their corresponding logs.



Figure 5. Guardian 360 Report cover page

Table of contents

1. Introduction	3
2. Scan Overview	4
3. Vulnerability Descriptions	6
Issue: [REDACTED]	7
Issue: [REDACTED]	8
Issue: [REDACTED]	9
Issue: [REDACTED]	10
Issue: [REDACTED]	11
Issue: [REDACTED]	12
Issue: [REDACTED]	13

Figure 6. Guardian 360 Report table of contents

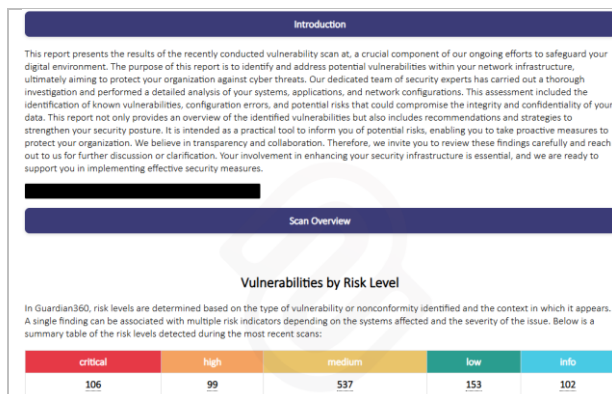


Figure 7. Guardian 360 Report Intro and risk table




Figure 8. Guardian 360 Report Donut chart




Figure 9. Guardian 360 Report stacked barchart

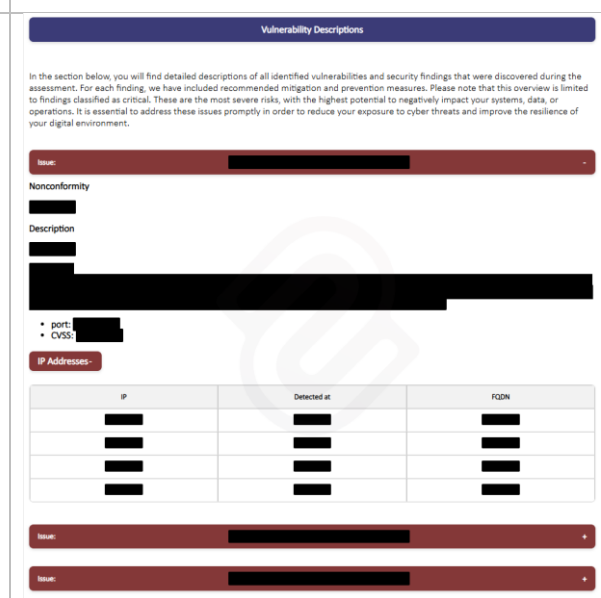


Figure 10. Issues Guardian 360 Issues pages


```

guardian360 > new > backups > 23-04 1653 issue page numbers, toggle section > gd_generate_report > ...
663   output_filename = output_dir / f'[{company}] - {scanobject} report [{lang}].html'
664
665   output_dir.mkdir(parents=True, exist_ok=True)
666
667   # write the file
668   with open(output_filename, 'w', encoding='utf-8') as f:
669       f.write(rendered_html)
670
671   # generate pdf version of the same report using Playwright
672   html_path_str = str(output_filename.resolve()) # full path to HTML
673   pdf_path = output_filename.with_suffix('.pdf') # same name, but with .pdf extension
674   # Run async Playwright logic to generate the PDF
675   asyncio.run(html_to_pdf(html_path_str, str(pdf_path)))
676
677   print(f'Saved HTML and PDF reports of: [{company}] - {scanobject} report [{lang}]')
678
679
680
681 # =====
682 # Debug output (Company-level row counts)
683 # =====
684
685 scanobject_risk_counts = defaultdict(lambda: defaultdict(int)) # Track counts by scanobject & risk
686 print(f'\nCompany Row Counts:')
687 for company, count in company_counts.items():
688     print(f'    {company}: {count} rows')
689
690 print('\nReports generated.')
691 print('Brought to life by Balsangur Dudayev, and bestowed upon (insert name) for worthy use.')
692
693
694 |
695 # <!-- <p class="client-title">{{ translations[lang]['Client Title'] }}</p>
696 # <p class="client-title">{{ translations[lang]['Scanobject Title'] }}</p>
697 # <p class="client-title">{{ translations[lang]['Scanobject Title'] }}</p>
698 # <p class="client-title">{{ translations[lang]['Scanobject Title'] }}</p>

```

Figure 11. Script used to generate report, 700 lines of code

```

guardian360 > new > backups > 23-04 1653 issue page numbers, toggle section > gd_reporttemplate.html > html > body
2   <html lang="en">
580   <body style="margin: 0; padding: 0;font-family:calibri light ;">
622   <div style="width: 1024px; margin: 0px auto; background-color: #ffff; padding: 10px; b
990   <script>
1066       function insertPageNumbersForTOC() {
1079           // Sort offsets just in case DOM order isn't perfect
1080           breakOffsets.sort((a, b) => a - b);
1081
1082           links.forEach(link => {
1083               const targetId = link.getAttribute('href').substring(1); // remove '#'
1084               const targetEl = document.getElementById(targetId);
1085
1086               if (targetEl) {
1087                   const rect = targetEl.getBoundingClientRect();
1088                   const scrollY = window.scrollY || document.documentElement.scrollTop;
1089                   const yOffset = rect.top + scrollY;
1090
1091                   // Count how many breaks occurred before this section
1092                   const pageIndex = breakOffsets.filter(offset => offset < yOffset).length;
1093                   const pageNumber = pageIndex + 2;
1094
1095                   const pageNumberSpan = document.createElement('span');
1096                   pageNumberSpan.textContent = pageNumber;
1097                   pageNumberSpan.style.float = 'right';
1098                   pageNumberSpan.style.marginLeft = '50px';
1099
1100                   pageNumberSpan.style.color = '#000';
1101                   pageNumberSpan.style.fontSize = '1em';
1102
1103                   link.appendChild(pageNumberSpan);
1104               }
1105           });
1106       }
1107
1108       window.addEventListener('load', insertPageNumbersForTOC);
1109
1110
1111   </script>
1112
1113
1114   </body>
1115 </html>
1116
1117

```

Figure 12. Script template used to generate report, 1100 lines of code

3.3. Integrating SIEM automation

I set up an n8n Docker container and configured a Nginx container as a reverse proxy. (Nginx, 2025)

I secured them with HTTPS using a mix of self-made certificates and ones that were already set up from Elasticsearch. I also created an IaC playbook for this setup, so it can be easily replicated on the live server after testing. Ensuring that n8n was up and running was crucial during this phase, as the work of all team members depended on it. (n8n, n.d.)

I built an automated workflow that pulls emails from a Gmail mailbox I set up. This automation looks at the content of the emails to decide whether to stop or continue contacting certain employees in the company through Vocal Notify. Vocal Notify is a tool that can be configured to call employees when specific alerts or events occur. However, if someone wants to stop receiving these calls, they can simply send an email to the mailbox I created.

To prevent the mailbox from requiring manual maintenance, I developed a Google Apps Script that deletes emails that have already been read, scheduled to run monthly. (Apps Script | Google for Developers, 2024)

The n8n workflow is configured to retrieve only unread emails and marks them as read after processing to avoid duplicate handling.

The original workflow using the Google Script was scrapped due to certain requests from upper management. Later, I changed and extended the workflow so that it could send out emails based on both the **type** and **volume** of alerts for Cyberalarm.




Figure 13. Running n8n container with nodes to and act on retrieve mails

Code

Test step

Parameters

Settings

Docs

Mode

Run Once for All Items

Language

JavaScript

JavaScript

```

1  const subject = $json["subject"];
2
3  // Match for the subject starting with "stop" or "start", followed by a first name
   and last name.
4  const match = subject.match(/(stop|start)\s+(\w+)\s+(\w+)/);
5
6  if (match) {
7      const action = match[1]; // 'stop' or 'start'
8      const firstName = match[2]; // First name
9      const lastName = match[3]; // Last name
10
11     return [
12         {
13             json: {
14                 action: action, // 'stop' or 'start'
15                 employee: `${firstName} ${lastName}`, // Full name
16             },
17         },
18     ];
19 } else {
20     return [{ json: { error: "Invalid subject format" } }];
21 }

```

Figure 14. Nn8n node code to take action on email's contents

Apps Script

Delete read emails

Implementeren

Bestanden

Code.gs

Bibliotheken

Services

1

function deleteReadEmails() {

2

// Search for read emails

3

var threads = GmailApp.search('is:read'); // This searches for all read emails

4

5

// Iterate through the threads and delete them

6

for (var i = 0; i < threads.length; i++) {

7

threads[i].moveToTrash();

8

}

9

}

Figure 15. Google app script to delete read mails

```

! n8n.yml
1 - name: Deploy n8n with Nginx and SSL
5 tasks:
77 - name: Stop and remove Nginx container
78   community.docker.docker_container:
79     name: nginx
80     state: absent
81
82 - name: Run n8n container
83   community.docker.docker_container:
84     name: n8n
85     image:
86     state: started
87     restart_policy: always
88     networks:
89     - name:
90   volumes:
91   -
92   -
93   env:
94     GENERIC_TIMEZONE: "Europe/Brussels"
95     TZ: "Europe/Brussels"
96     NODE_EXTRA_CA_CERTS:
97
98 - name: Run Nginx container
99   community.docker.docker_container:
100     name: nginx
101     image:
102     state: started
103     restart_policy: always
104     published_ports:
105     -
106   volumes:
107   -
108   -
109   networks:
110   - name:
111

```

Figure 16. Playbook to run n8n and NGINX container

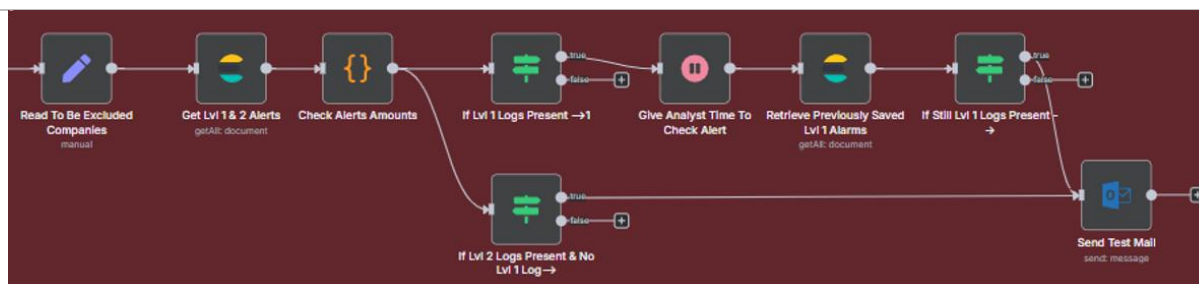


Figure 17. Updated Vocal Notify workflow to make calls based on logs

3.4. Integrating TheHive

I conducted research on both TheHive and MISP, including completing their respective TryHackMe rooms. Based on this analysis, I chose to proceed with installing TheHive and decided to omit MISP.

During the setup, I deployed the required Cassandra container alongside TheHive. TheHive also requires Nginx and Elasticsearch to function properly. Since we already had existing containers for both Nginx and Elasticsearch in our SOC environment, I configured TheHive to connect to those instead of creating new ones. This approach helped optimize resource usage and improved the overall efficiency of the system.

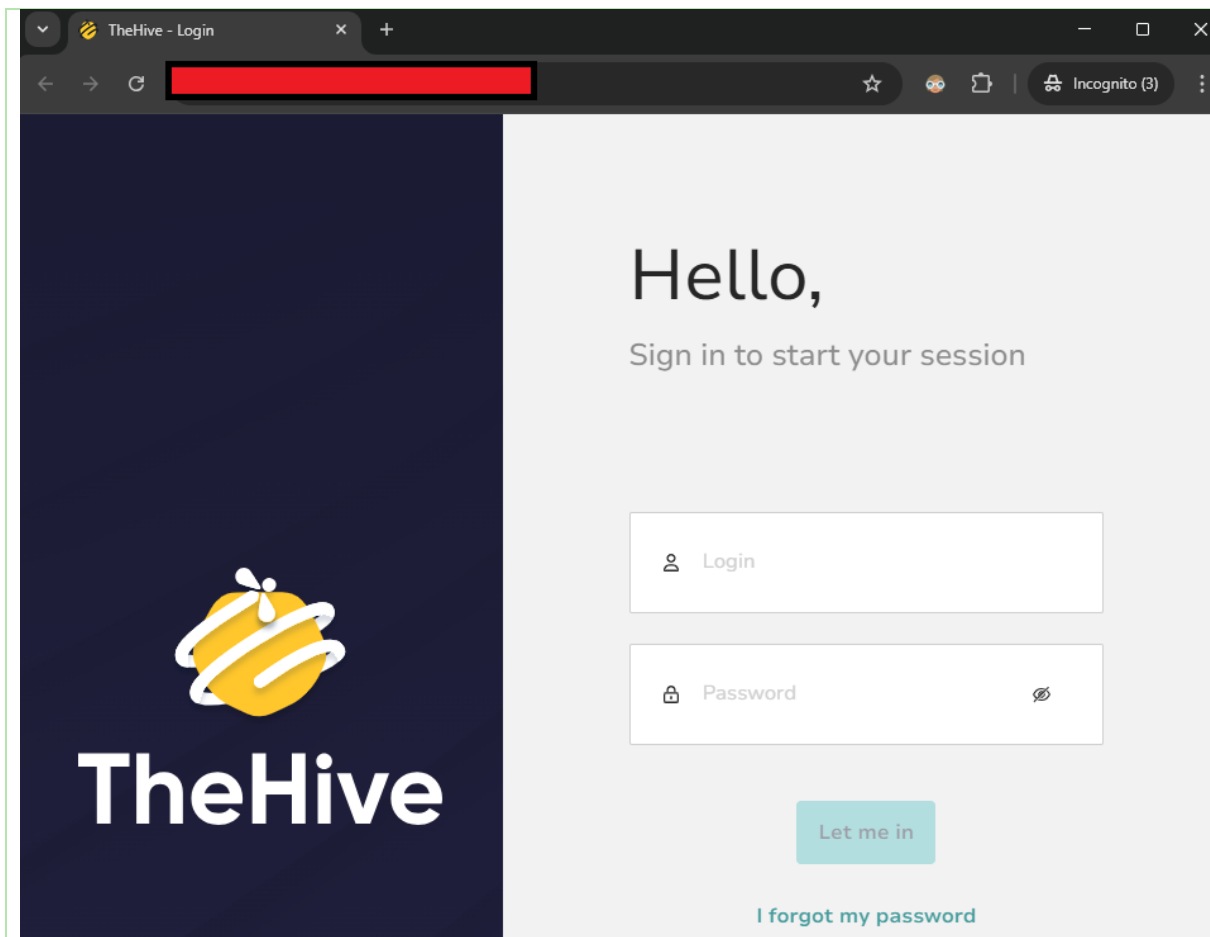


Figure 18. TheHive running

4. Conclusion

This report summarizes the progress made during the internship focused on developing and enhancing a fully functional Security Information and Event Management (SIEM) solution for Netcare. The SIEM environment aims to centralize log collection, alerting, and reporting to improve security monitoring both for the company and its clients.

The project followed a structured timeline, starting with the deployment of the ELK stack components and automated infrastructure using Ansible to enable reliable, real-time log collection and visualization. A multilingual reporting system was developed to transform vulnerability scan data into interactive reports, supporting informed decision-making.

The integration of n8n facilitated workflow automation and improved alert management by connecting with a custom mailbox, causing Vocal Notify alerts to be more efficient.

The SOC environment consists of an intelligent, automated platform that consolidates diverse security data sources, improving situational awareness and operational efficiency.

5. References

- Apps Script* | *Google for Developers*. (2024, 31 10). Opgeroepen op 4 2025, van developers.google.com: <https://developers.google.com/apps-script>
- Deepl. (2025, 4 23). *Introduction* | *DeepL API Documentation*. Opgeroepen op 4 23, 2025, van developers.deepl.com: <https://developers.deepl.com/docs>
- Elastic. (2025, 4 23). *Logstash* | *Elastic Documentation*. Opgehaald van elastic.co: <https://www.elastic.co/docs/reference/logstash>
- elasticsearch. (2025, 04 23). *Index basics* | *Elastic Docs*. Opgehaald van www.elastic.co: <https://www.elastic.co/docs/manage-data/data-store/index-basics>
- n8n. (sd). *Docker* | *n8n Docs*. Opgeroepen op 04 23, 2025, van docs.n8n.io: <https://docs.n8n.io/hosting/installation/docker/>
- Nginx. (2025, 3 23). *NGINX Reverse Proxy* | *NGINX Documentation*. Opgehaald van <https://docs.nginx.com/nginx/>: <https://docs.nginx.com/nginx/admin-guide/web-server/reverse-proxy/>
- Stringer, T. (2025, 4 3). *Developing modules* | *Ansible Community Documentation*. Opgeroepen op 4 3, 2025, van docs.ansible.com: https://docs.ansible.com/ansible/latest/dev_guide/developing_modules_general.html#environment-setup